

all librosa code and feature explanation from the [CCRMA Music Information Retrieval summer intensive github repo \(https://github.com/bmcfee/stanford-mir\)](https://github.com/bmcfee/stanford-mir)

```
In [104]: %matplotlib inline
import numpy, scipy, matplotlib.pyplot as plt, IPython.display as ipd
import librosa, librosa.display
import seaborn as sns
import sklearn
import pandas as pd
plt.rcParams['figure.figsize'] = (14, 5)
```

```
In [295]: pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
```

```
In [2]: sns.set_style('whitegrid')
```

load audio

load the audio from the "forest" video:

```
In [3]: sns.set_style('whitegrid')
```

```
In [4]: x, sr = librosa.load('../..../audioFiles/forest.wav')
```

```
In [5]: ipd.Audio(x, rate=sr)
```

Out[5]:

0:00 / 2:51

define helper functions to normalize

```
In [205]: # normalize data between 0 and 1 for visualization
def normalize(x, axis=0):
    return sklearn.preprocessing.minmax_scale(x, axis=axis)
```

```
In [206]: # normalize data to mean of 0 and unit variance
```

extract root mean square energy

```
In [224]: hop_length = 512
         frame_length = 1024
```

```
In [225]: energy = numpy.array([
            sum(abs(x[i:i+frame_length]**2))
            for i in range(0, len(x), hop_length)
        ])
```

```
In [226]: energy.shape
```

```
Out[226]: (7383,)
```


```
In [227]: rms = librosa.feature.rms(x, frame_length=frame_length, hop_length=hop_length)[0]
```

```
In [228]: rms.shape
```

```
Out[228]: (7383,)
```

```
In [229]: frames = range(len(energy))
         t = librosa.frames_to_time(frames, sr=sr, hop_length=hop_length)
```

```
In [230]: librosa.display.waveplot(x, sr=sr, alpha=0.4)
         plt.plot(t, energy/energy.max(), 'r--') # normalized for visualization
         plt.plot(t[:len(rms)], rms/rms.max(), color='g') # normalized for visualization
         plt.legend(('Energy', 'RMSE'))
         plt.show()
```



```
In [232]: rms = sklearn.preprocessing.scale(rms)
```

```
In [233]: rms.mean()
```

```
Out[233]: -8.266986e-09
```

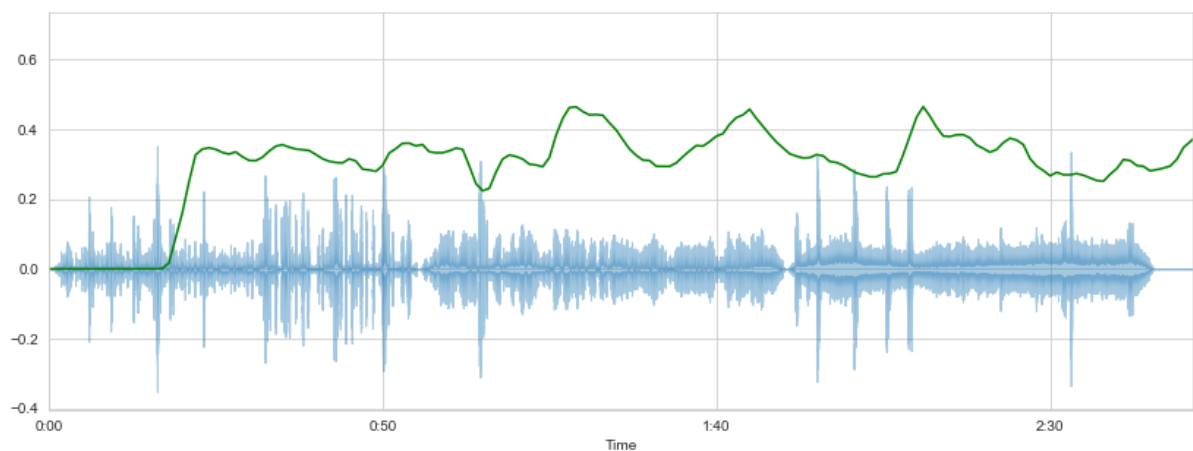
```
In [235]: rms.var()
```

```
Out[235]: 0.99999994
```

extract zero line crossing rate

Add a small constant to avoid oscillation around silence triggering high zero-crossing rate:

```
In [243]: librosa.display.waveplot(x, sr=sr, alpha=0.4)
          zcrs = librosa.feature.zero_crossing_rate(x + 0.0001)
          plt.plot(zcrs[0], 'g')
          plt.show()
```



```
In [244]: zcrs.shape
```

```
Out[244]: (1, 7383)
```

extract spectral features

extract spectral centroid

The **spectral centroid** ([Wikipedia \(https://en.wikipedia.org/wiki/Spectral_centroid\)](https://en.wikipedia.org/wiki/Spectral_centroid)) indicates at which frequency the energy of a spectrum is centered upon. This is like a weighted mean:


$$f_c = \frac{\sum_k S(k) f(k)}{\sum_k S(k)}$$

where $S(k)$ is the spectral magnitude at frequency bin k , $f(k)$ is the frequency at bin k .

`librosa.feature.spectral_centroid`
https://librosa.github.io/librosa/generated/librosa.feature.spectral_centroid.html#librosa.feature.spectral_centroid
 computes the spectral centroid for each frame in a signal:

```
In [198]: # calculate time variable for plotting
frames = range(len(spectral_centroids))
t = librosa.frames_to_time(frames)
```

```
In [245]: spectral_centroids = librosa.feature.spectral_centroid(x+0.01, sr=sr)[0]
# add constant to correct value at silence
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_centroids), color='g') # normalize for visualization purposes
plt.show()
```



```
In [246]: spectral_centroids = sklearn.preprocessing.scale(spectral_centroids)
```

```
In [248]: spectral_centroids.mean()
```

```
Out[248]: 6.159384408593582e-17
```

```
In [249]: spectral_centroids.var()
```

```
Out[249]: 1.0000000000000002
```

extract spectral bandwidth

[librosa.feature.spectral_bandwidth](#)


(https://librosa.github.io/librosa/generated/librosa.feature.spectral_bandwidth.html#librosa.feature.spectral_bandwidth)

computes the order- p spectral bandwidth:

$$\left(\sum_k S(k) \left(f(k) - f_c \right)^p \right)^{\frac{1}{p}}$$

where $S(k)$ is the spectral magnitude at frequency bin k , $f(k)$ is the frequency at bin k , and f_c is the spectral centroid. When $p = 2$, this is like a weighted standard deviation.

```
In [250]: spectral_bandwidth_2 = librosa.feature.spectral_bandwidth(x+0.01, sr=sr)
[0]
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_bandwidth_2), color='g')
plt.show()
```



```
In [251]: spectral_bandwidth_2 = sklearn.preprocessing.scale(spectral_bandwidth_2)
```

```
In [252]: spectral_bandwidth_2.mean()
```

```
Out[252]: -1.5398461021483954e-16
```


```
In [253]: spectral_bandwidth_2.var()
```

```
Out[253]: 1.0000000000000002
```

extract spectral rolloff

Spectral rolloff is the frequency below which a specified percentage of the total spectral energy, e.g. 85%, lies.

```
In [254]: spectral_rolloff = librosa.feature.spectral_rolloff(x+0.01, sr=sr)[0]
librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_rolloff), color='g')
plt.show()
```



```
In [255]: spectral_rolloff = sklearn.preprocessing.scale(spectral_rolloff)
```

```
In [256]: spectral_rolloff.mean()
```

```
Out[256]: 1.3858614919335558e-16
```

```
In [257]: spectral_rolloff.var()
```

```
Out[257]: 0.9999999999999997
```


extract spectral novelty

We will compute a **spectral novelty function** (FMP (<https://www.audiolabs-erlangen.de/resources/MIR/FMP/C0/C0.html>), p. 309):

1. Compute the log-amplitude spectrogram.
2. Within each frequency bin, k , compute the energy novelty function as shown earlier, i.e. (a) first-order difference, and (b) half-wave rectification.
3. Sum across all frequency bins, k .

```
In [258]: spectral_novelty = librosa.onset.onset_strength(x, sr=sr)
```

```
In [259]: librosa.display.waveplot(x, sr=sr, alpha=0.4)
plt.plot(t, normalize(spectral_novelty), 'g')
plt.xlim(0, t.max())
plt.xlabel('Time (sec)')
plt.legend(('Spectral Novelty',))
plt.show()
```



```
In [260]: spectral_novelty = sklearn.preprocessing.scale(spectral_novelty)
```

```
/Users/trqk-data/.local/share/virtualenvs/nature-nurtures--FcyQ97q/lib/
python3.7/site-packages/sklearn/preprocessing/data.py:189: UserWarning:
Numerical issues were encountered when scaling the data and might not b
e solved. The standard deviation of the data is probably very close to
0.
```

```
warnings.warn("Numerical issues were encountered ")
```

```
In [261]: spectral_novelty.mean()
```

```
Out[261]: 0.0
```

```
In [262]: spectral_novelty.var()
```

```
Out[262]: 1.0000001
```

Extract Mel Frequency Cepstral Coefficients (MFCCs)

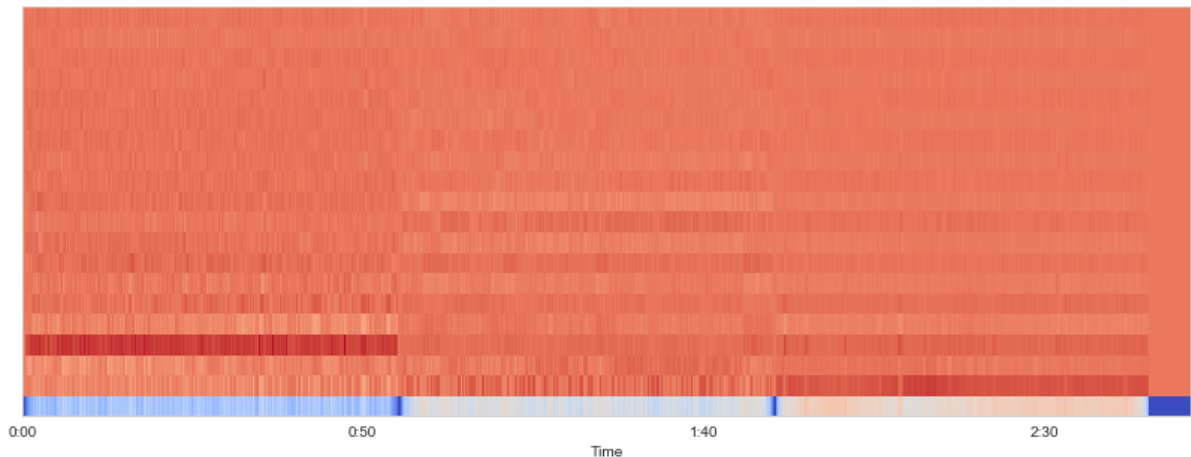
MFCCs

```
In [263]: mfccs = librosa.feature.mfcc(x, sr=sr)
print(mfccs.shape)
```

```
(20, 7383)
```

display MFCCs:

```
In [264]: librosa.display.specshow(mfccs, sr=sr, x_axis='time')
plt.show()
```



Let's scale the MFCCs such that each coefficient dimension has zero mean and unit variance:

```
In [265]: mfccs = sklearn.preprocessing.scale(mfccs, axis=1)
print(mfccs.mean(axis=1))
print(mfccs.var(axis=1))

[ 2.4478028e-08 -1.7526979e-08 -3.9790925e-07 -1.9843995e-07
 -1.8713743e-08 -4.1657859e-09 -1.2973678e-08 -2.5059300e-08
 2.6544775e-08 -3.5457617e-08 2.5083521e-08 3.5360741e-09
 3.8117747e-08 -4.5484569e-08 5.3156153e-08 2.2451648e-08
 -3.4957583e-08 -3.3407019e-08 -1.0148048e-08 5.8999152e-08]
[0.99999684 0.9999978 1.0000142 0.9999928 1.0000222 1.0000057
 1.0000012 1.0000045 0.99999297 0.999995 0.9999906 1.0000012
 1.0000141 0.9999916 0.9999927 1.0000049 0.99999785 1.0000064
 0.99999017 1.0000147 ]
```


```
/Users/trqk-data/.local/share/virtualenvs/nature-nurtures--FcyQ97q/lib/
python3.7/site-packages/sklearn/preprocessing/data.py:172: UserWarning:
Numerical issues were encountered when centering the data and might not
be solved. Dataset may contain too large values. You may need to presca
le your features.
```

```
warnings.warn("Numerical issues were encountered ")
/Users/trqk-data/.local/share/virtualenvs/nature-nurtures--FcyQ97q/lib/
python3.7/site-packages/sklearn/preprocessing/data.py:189: UserWarning:
Numerical issues were encountered when scaling the data and might not b
e solved. The standard deviation of the data is probably very close to
0.
```

```
warnings.warn("Numerical issues were encountered ")
```


Display all the scaled MFCCs:


```
In [266]: librosa.display.specshow(mfccs, sr=sr, x_axis='time')  
plt.show()
```



Pull out one MFCC:

```
In [267]: librosa.display.waveplot(x, sr=sr, alpha=0.4)  
plt.plot(t, mfccs[1], color='g')  
plt.show()
```



Create Feature Vector

Append spectral features to MFCCs as new rows of the array:

```
In [283]: def add_second_dimension_to_array(the_array):  
           cols = the_array.shape[0]  
           return numpy.reshape(the_array, (-1, 7383))
```

```
In [284]: mfccs.shape
```

```
Out[284]: (20, 7383)
```

```
In [285]: spectral_features = [spectral_centroids,
                             spectral_bandwidth_2,
                             spectral_rolloff,
                             spectral_novelty,
                             ]
```

```
In [286]: X = mfccs
for feature in spectral_features:
    X = numpy.append(X, add_second_dimension_to_array(feature), axis=0)
```

```
In [287]: X.shape
```

```
Out[287]: (24, 7383)
```

Apply Principal Component Analysis to See Which Features Predict the Signal Most

```
In [288]: X.mean()
```

```
Out[288]: -4.024665683990886e-09
```

```
In [289]: model = sklearn.decomposition.PCA(n_components=2, whiten=True)
```

```
In [290]: model.fit(X.T)
```

```
Out[290]: PCA(copy=True, iterated_power='auto', n_components=2, random_state=None,
              svd_solver='auto', tol=0.0, whiten=True)
```

```
In [291]: Y = model.transform(X.T)
```

```
In [292]: print(Y.shape)
```

```
(7383, 2)
```

```
In [293]: model.components_.shape
```

```
Out[293]: (2, 24)
```

```
In [296]: pd.DataFrame(model.components_, index = ['PC-1', 'PC-2'])
```

```
Out[296]:
```

	0	1	2	3	4	5	6	7	
PC-1	-0.260838	-0.208214	-0.267103	0.315148	-0.163073	0.004098	0.106031	0.034172	0.30631
PC-2	0.163501	-0.296752	0.049073	0.243706	-0.294038	0.326705	-0.173676	0.278562	-0.01521

Component 1 depends substantially on MFCCs 4, 9, and 11, while component 2 depends largely on MFCCs 4-6, as well as spectral rolloff and spectral novelty.

Apply Non-Negative Matrix Factorization to Describe Signal as a Sum of Components


```
In [297]: S = librosa.stft(x)
print(S.shape)

(1025, 7383)
```

```
In [298]: Smag = librosa.amplitude_to_db(S)
librosa.display.specshow(Smag, sr=sr, x_axis='time', y_axis='log')
plt.colorbar()
plt.show()
```

/Users/trqk-data/.local/share/virtualenvs/nature-nurtures--FcyQ97q/lib/python3.7/site-packages/librosa/core/spectrum.py:1700: UserWarning: amplitude_to_db was called on complex input so phase information will be discarded. To suppress this warning, call amplitude_to_db(np.abs(S)) instead.

warnings.warn('amplitude_to_db was called on complex input so phase '



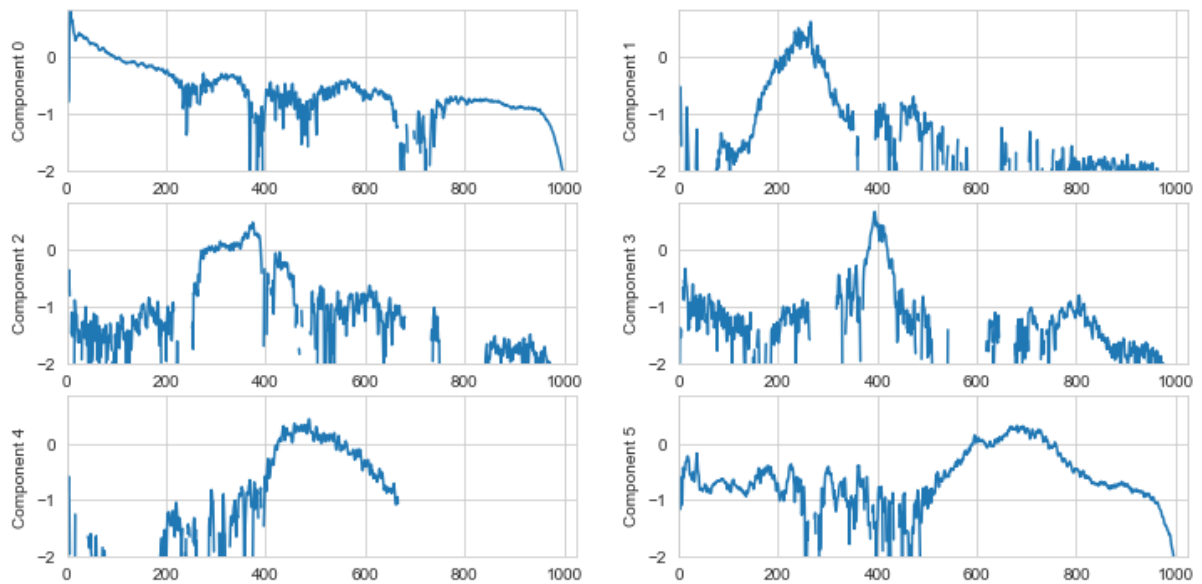
```
In [113]: X = numpy.absolute(S)
n_components = 6
W, H = librosa.decompose.decompose(X, n_components=n_components, sort=True)
print(W.shape)
print(H.shape)

(1025, 6)
(6, 7383)
```

Plot each spectral component:

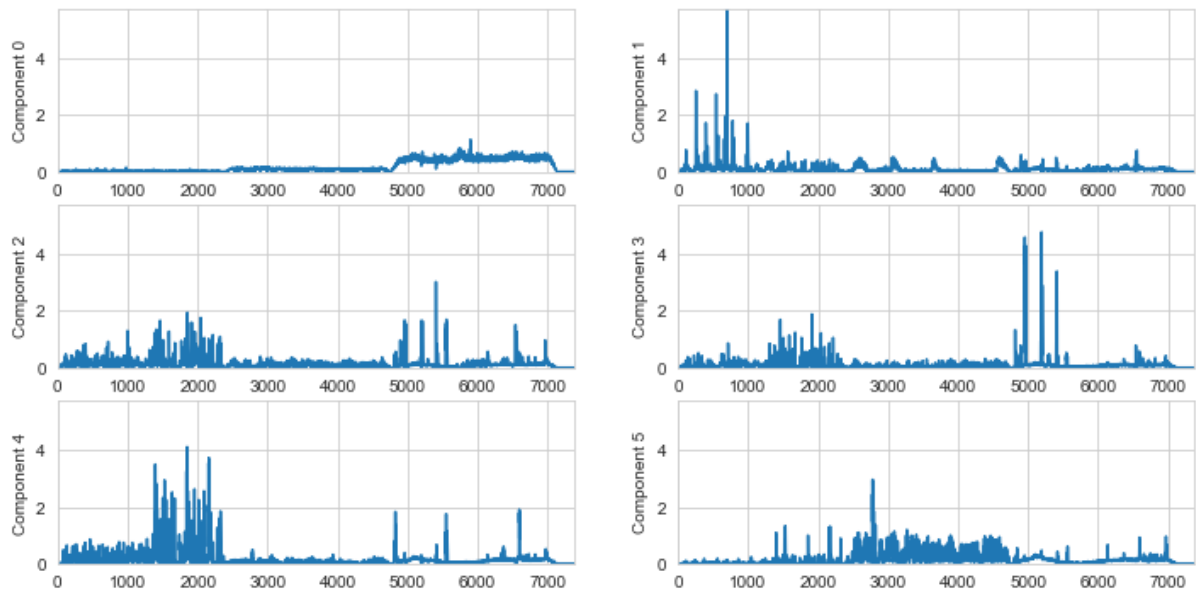
```
In [114]: plt.figure(figsize=(12, 6))
logW = numpy.log10(W)
for n in range(n_components):
    plt.subplot(numpy.ceil(n_components/2.0), 2, n+1)
    plt.plot(logW[:,n])
    plt.ylim(-2, logW.max())
    plt.xlim(0, W.shape[0])
    plt.ylabel('Component %d' % n)
```

/Users/trqk-data/.local/share/virtualenvs/nature-nurtures--FcyQ97q/lib/python3.7/site-packages/ipykernel_launcher.py:2: RuntimeWarning: divide by zero encountered in log10



temporal activations of each component:

```
In [115]: plt.figure(figsize=(12, 6))  
for n in range(n_components):  
    plt.subplot(numpy.ceil(n_components/2.0), 2, n+1)  
    plt.plot(H[n])  
    plt.ylim(0, H.max())  
    plt.xlim(0, H.shape[1])  
    plt.ylabel('Component %d' % n)
```



Listen to components:

```
In [116]: reconstructed_signal = scipy.zeros(len(x))
          for n in range(n_components):
              Y = scipy.outer(W[:,n], H[n])*numpy.exp(1j*numpy.angle(S))
              y = librosa.istft(Y)
              reconstructed_signal[:len(y)] += y
          ipd.display( ipd.Audio(y, rate=sr) )
```

0:00 / 2:51

0:00 / 2:51

0:00 / 2:51

0:00 / 2:51

0:00 / 2:51

0:00 / 0:00